PART II
ANALYTICS

Microsoft BI Timeline

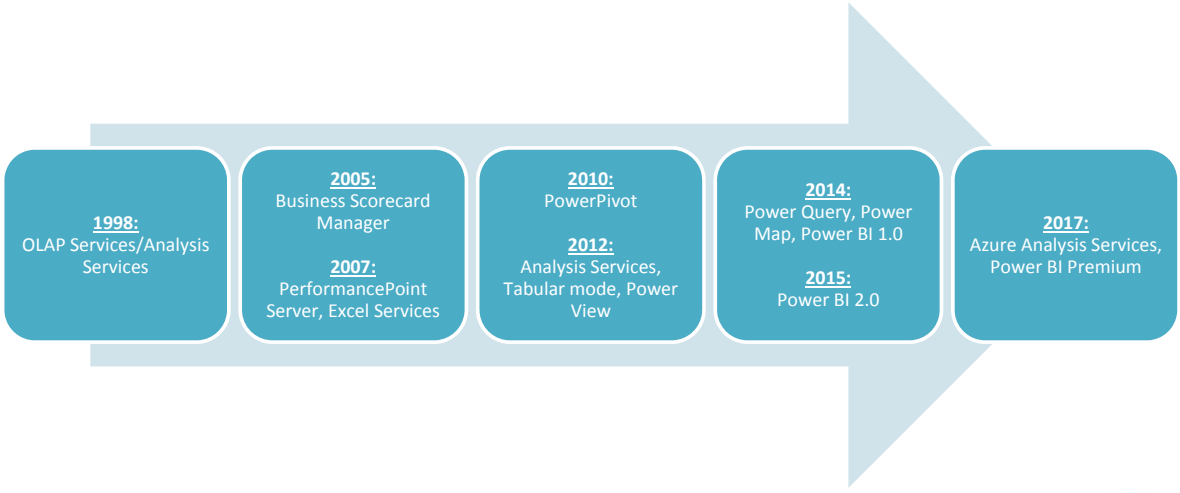

Microsoft Big Data Timeline

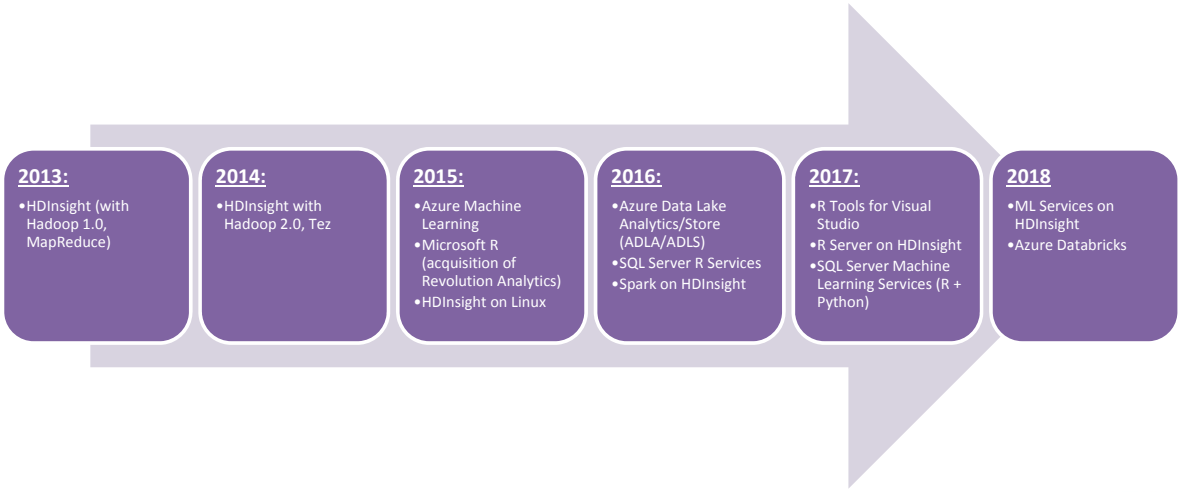

SQL Server Analytics Timeline

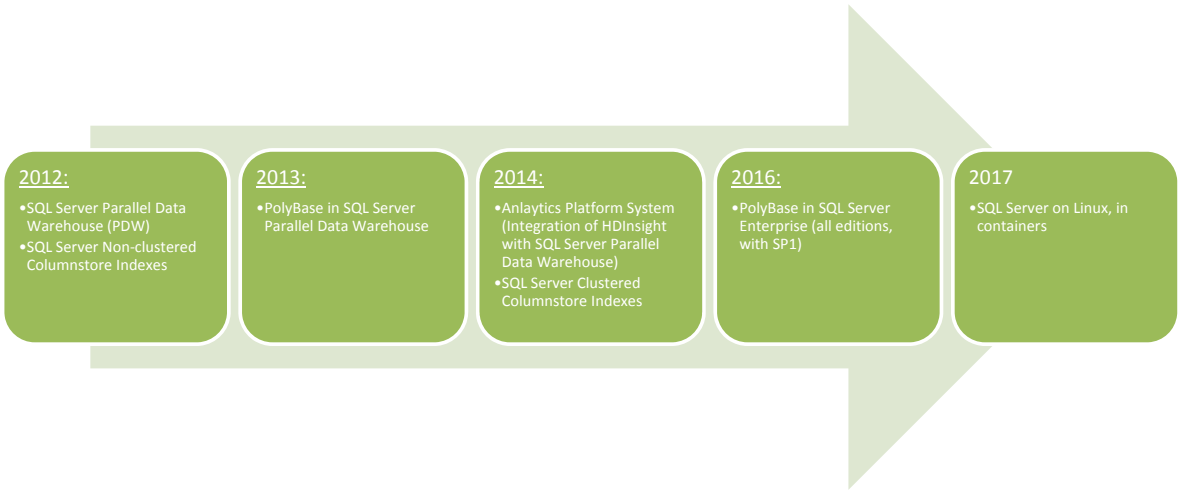

The Result: Lots to Learn

- There are so many components
- Each is rich and complex
- But they all connect
- Let's look at ways to slice this…

Break it down by: Repositories

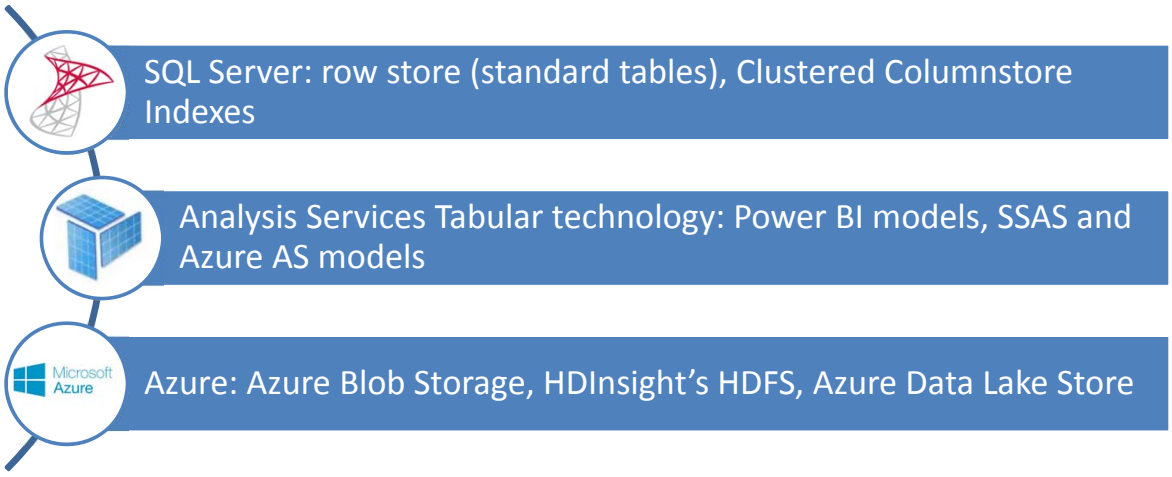

Break it down by: Query front-ends

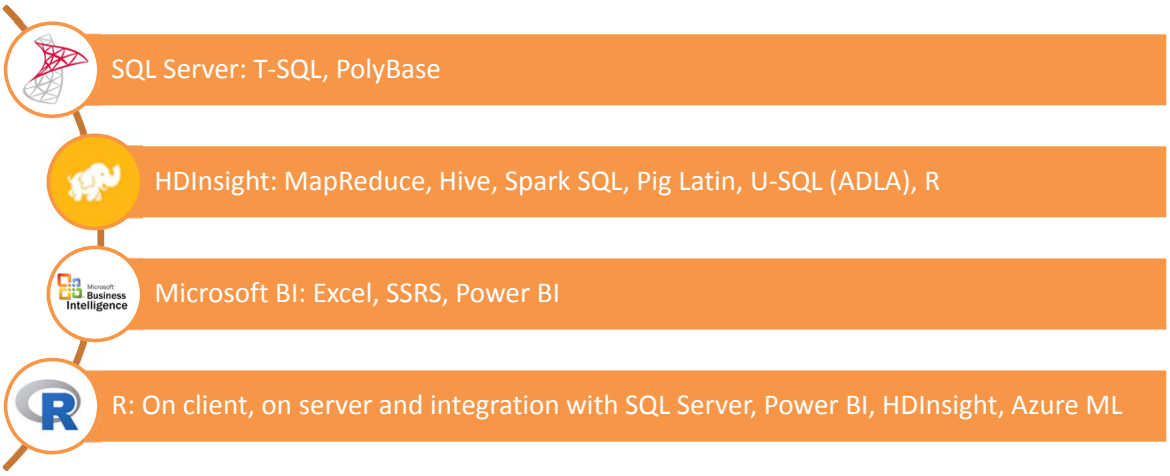

Break it down by: Client Tools

How to think about it

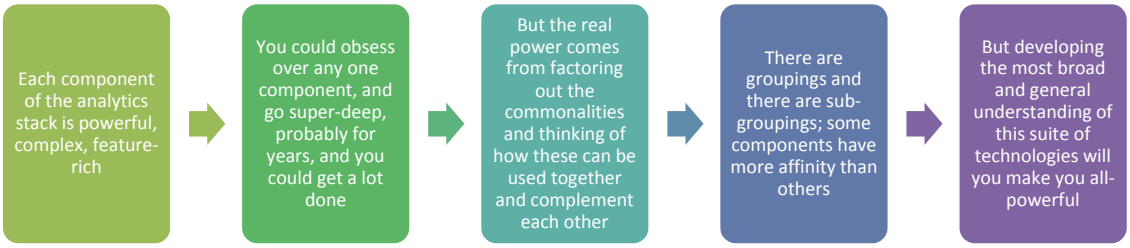

In other words…

Our core agenda, stretch goals and exceptions…

WHAT WE'LL COVER

Our Agenda: Preliminaries

- Data warehouse concepts
- Introducing our dataset (NYC 311 service calls)
- Open in Power BI, inspect

Our Agenda: Big Data

- Discuss HDInsight
- Query and process in Hadoop
  - MapReduce (separate data set)
  - Hive and Pig
- Further process in Azure Data Lake Analytics/U-SQL
- SQL Server PolyBase and Clustered Columnstore Indexes
- Apache Spark

Our Agenda: Business Intelligence

- Power BI – deeper dive
- Analysis Services
- Azure Analysis Services

Dimensional Analysis and MPP

DATA WAREHOUSE CONCEPTS

Business Intelligence

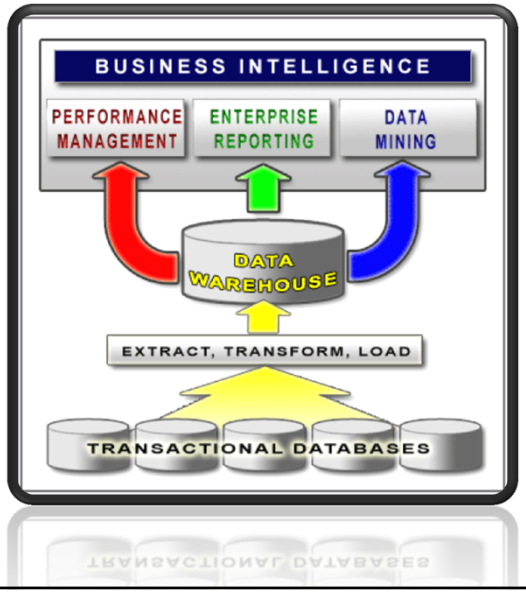

Preparing For Business Intelligence

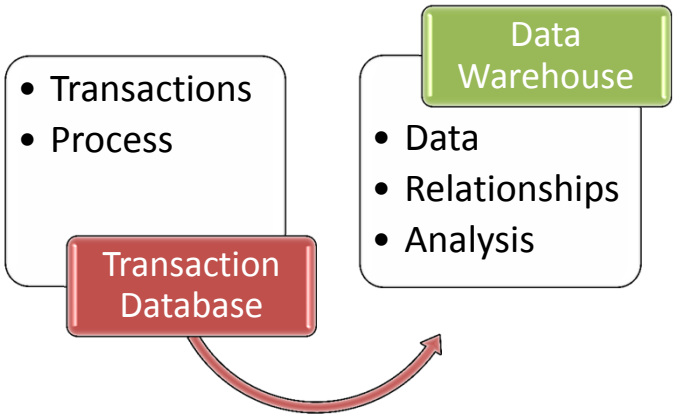

Dimensional Model

- Meas
- Dime
- Hiera
- Grain
- Star S

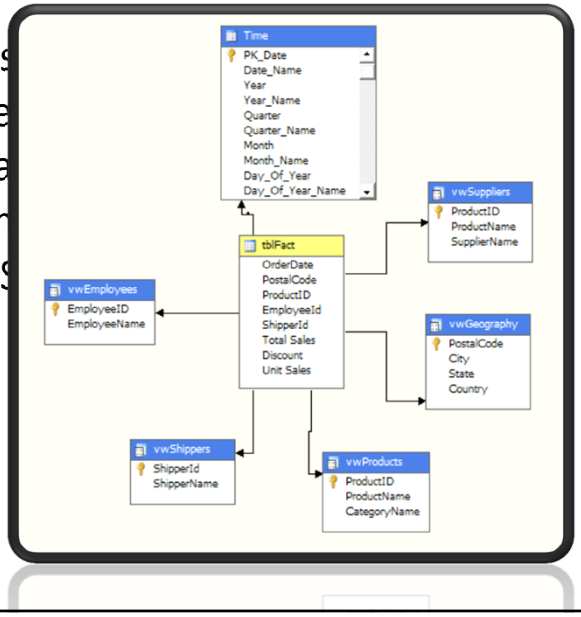

Star Schemas

- Physical data model
- Central fact table
- Multiple dimension tables
  - Used to constrain fact table queries

Example Data Request

- Get Total Sales By State, By Month for a Calendar Year For Country = USA and Calendar Year = 1996

Data Warehouse Query

	STATE	Month_Name	(No column name)
1	NM	August 1996	3343.60
2	WY	August 1996	48.00
3	ID	December 1996	6038.60
4	OR	December 1996	780.00
5	WY	December 1996	3391.20
6	NM	July 1996	624.80
7	WA	July 1996	676.00
8	NM	November 1996	1731.20
9	WA	November 1996	2856.00
10	WY	November 1996	141.60
11	AK	October 1996	934.50

What is MPP?

- Massively Parallel Processing
  - A cluster of individual RDBMS instances (worker nodes)
  - One master node, in front
    - Takes query, delegates parts of it to different worker nodes
    - Combines worker nodes' results, returns as single result set
  - Thus, appears as a single RDBMS
    - Send it one query, get back one result set
    - But query is highly parallelized, so it's fast
    - Perfect for data warehouses
  - Bears some resemblance to MapReduce
  - Examples include Teradata, HP Vertica, IBM Netezza, Pivotal Greenplum

What is MPP?

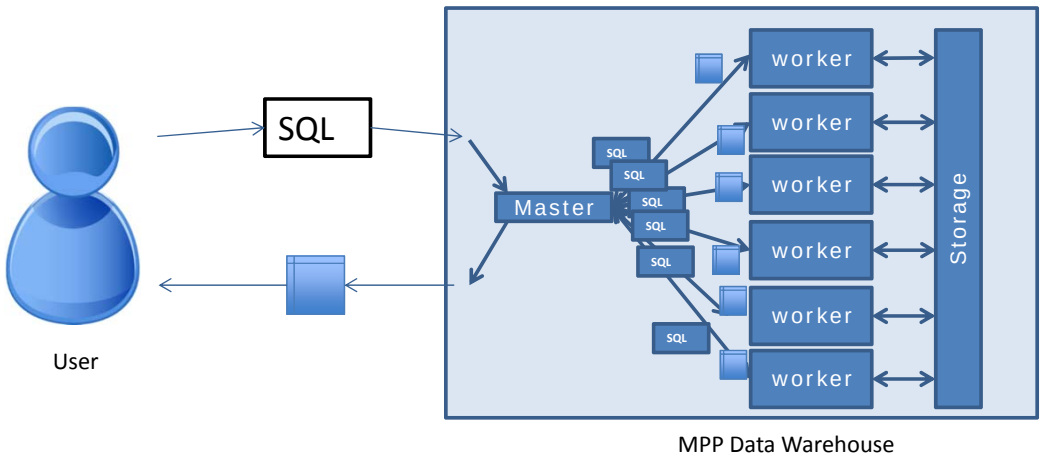

Column-Oriented Stores

- Imagine, instead of:

Employee ID	Age	Income
1	43	90000
2	38	100000
3	35	100000

- You have:

Employee ID	1	2	3
Age	43	38	35
Income	90000	100000	100000

- Perf: values you wish to aggregate are adjacent
- Efficiency: great compression from identical or nearly-identical values in proximity
- Fast aggregation and high compression means huge volumes of data can be stored and processed, in RAM

MPP + Columnar

- Together, these greatly accelerate DW performance.
- Far superior to a scaled-up SQL Server Enterprise box
- Most DW platforms combine these two technologies
- Add vector processing and it's a big deal

MPP at Microsoft?

- Yes, resulting from 2008 acquisition of DATAllegro
  - Open source MPP based on Ingres, written in Java, running on Linux
- Project Madison
  - Apply DATAllegro architecture using SQL Server, .NET and Windows
  - Released as SQL Server Parallel Data Warehouse (PDW)
  - Now called Analytics Platform System (APS)

NYC 311 Service Calls

OUR DATA SET

NYC Open Data

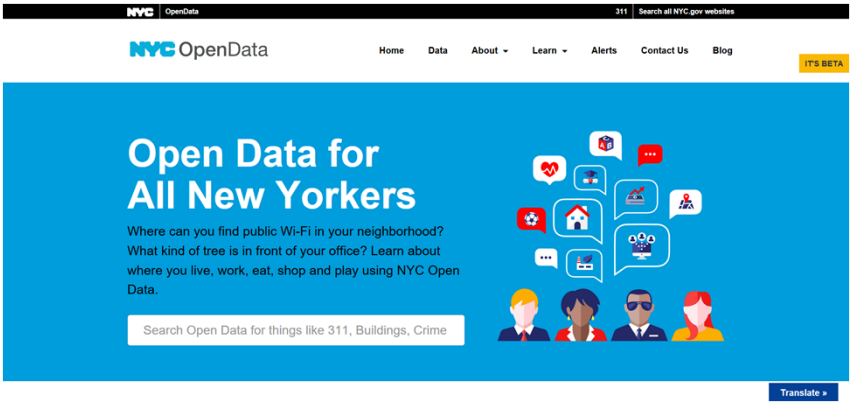

- https://data.cityofnewyork.us/Social-Services/311-Service-Requests-from-2010-to-Present/erm2-nwe9

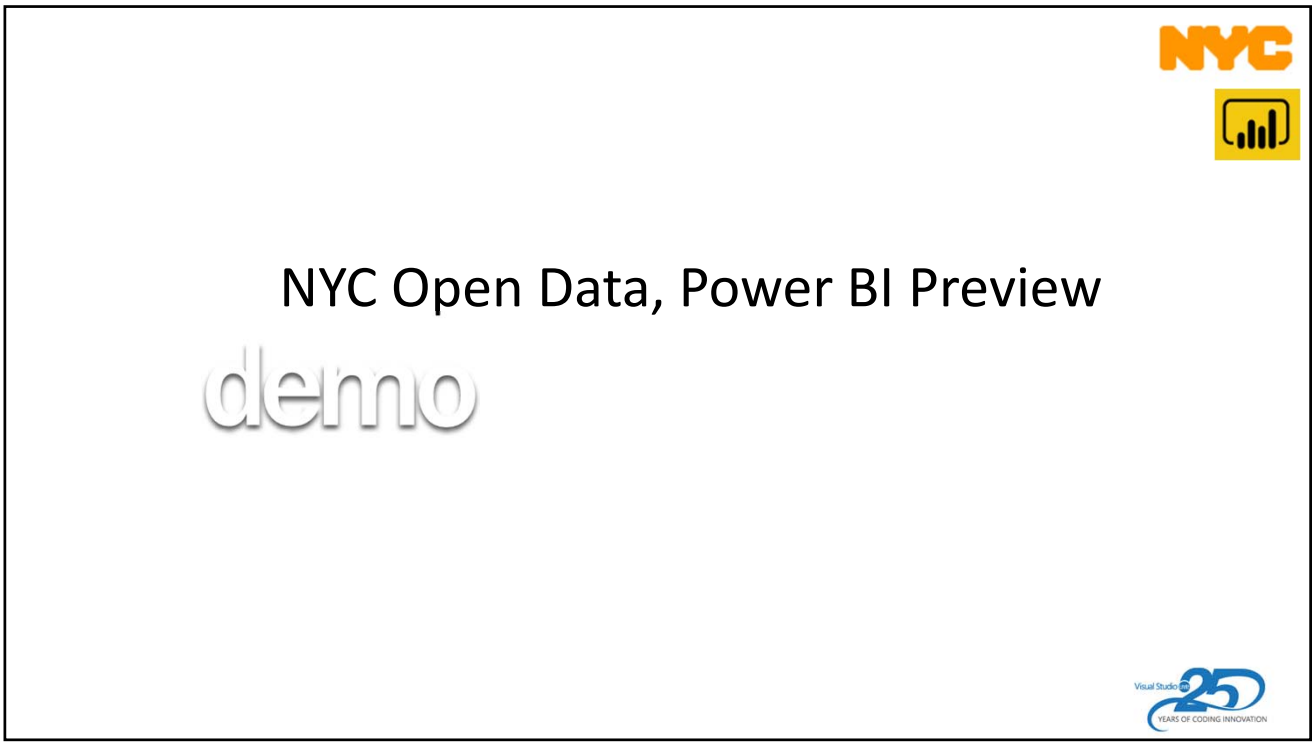
NYC Open Data, Power BI Preview
demo

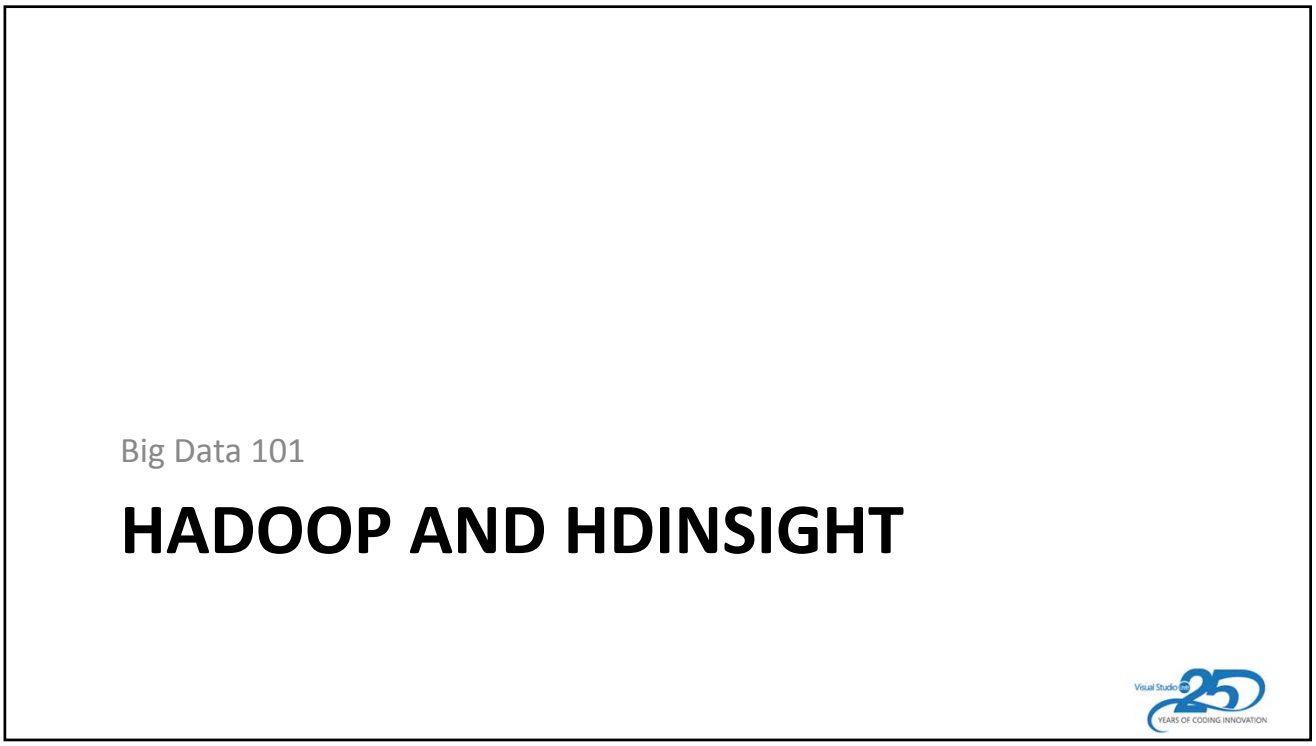
Big Data 101
HADOOP AND HDINSIGHT

What is Big Data?

- 100s of TB into PB and higher
- Involving data from: financial data, sensors, web logs, social media, etc.
- Parallel processing often involved
  - Hadoop is emblematic, but other technologies are Big Data too
- Processing of data sets too large for transactional databases
  - Analyzing *interactions*, rather than *transactions*
  - The three V's: Volume, Velocity, Variety
- Big Data tech sometimes imposed on small data problems

What's a "Data Lake?"

- Definition #1: The Big Data version of a data warehouse
- Definition #2: A place where you land all the data you don't know what to do with (aka a Data "Swamp")
- Definition #3: A file system repository where raw data is stored in file form (formats ranging from CSV to JSON to Hadoop sequence files to Apache Parquet)
  - HDFS, Amazon S3, Azure BLOB storage, Azure Data Lake Store
- Definition #4: A set of technologies that treat files or folders like (big) tables

What's MapReduce?

- "Big" data input accepted in file form
- Data is partitioned and sent to *mappers* (nodes in cluster)
- Mappers pre-process data into KV pairs, then all output for (a) given key(s) goes to a *reducer*
- Reducers aggregate; one line of output per unique key, with one value
- Map and Reduce code natively written as Java functions

MapReduce, in a Diagram

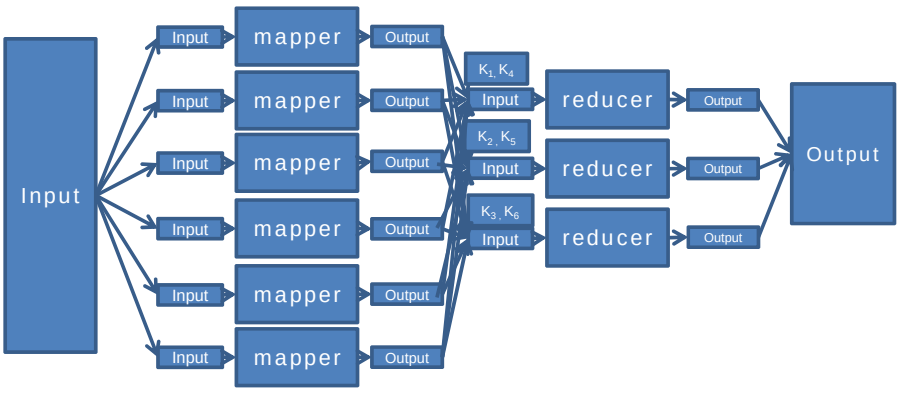

Apache Tez

- Key component added to Hadoop 2.0
- It's a directed acyclic graph (DAG) execution engine that runs on top of YARN (Hadoop 2.0's resource manager)
- Hive and Pig can both run on it
- Shunned by Cloudera

HDFS

- File system whose data gets distributed over commodity drives on commodity servers
- Data is replicated
- If one box goes down, no data lost
  - "Shared Nothing"
  - Except the name node
- BUT: Immutable
  - Files can only be written to once
  - So updates require drop + re-write (slow)
  - You can append though
  - Like a DVD/CD-ROM

Hadoop 3,
Open Hybrid Architecture Initiative

- Hadoop 3: YARN jobs as Docker containers
- Open Hybrid Architecture Initiative
  - Separate storage from compute
    - Ozone file system sub-project
  - Containerize Hadoop -- deploy to Kubernetes clusters
  - Will allow Hadoop environments to move between on-prem and cloud; and/or across multiple clouds
  - This is just starting

HDINSIGHT

Microsoft HDInsight

- Developed with Hortonworks and incorporates Hortonworks Data Platform (HDP) *for Windows*
- Windows Azure HDInsight and Microsoft HDInsight Server
  - Single node preview runs on Windows client
  - Also Hortonworks HDP for Windows
  - Also HDInsight with Analytics Platform System
- Includes ODBC Drivers for Hive
- All contributed back to open source Apache project

Azure HDInsight Provisioning

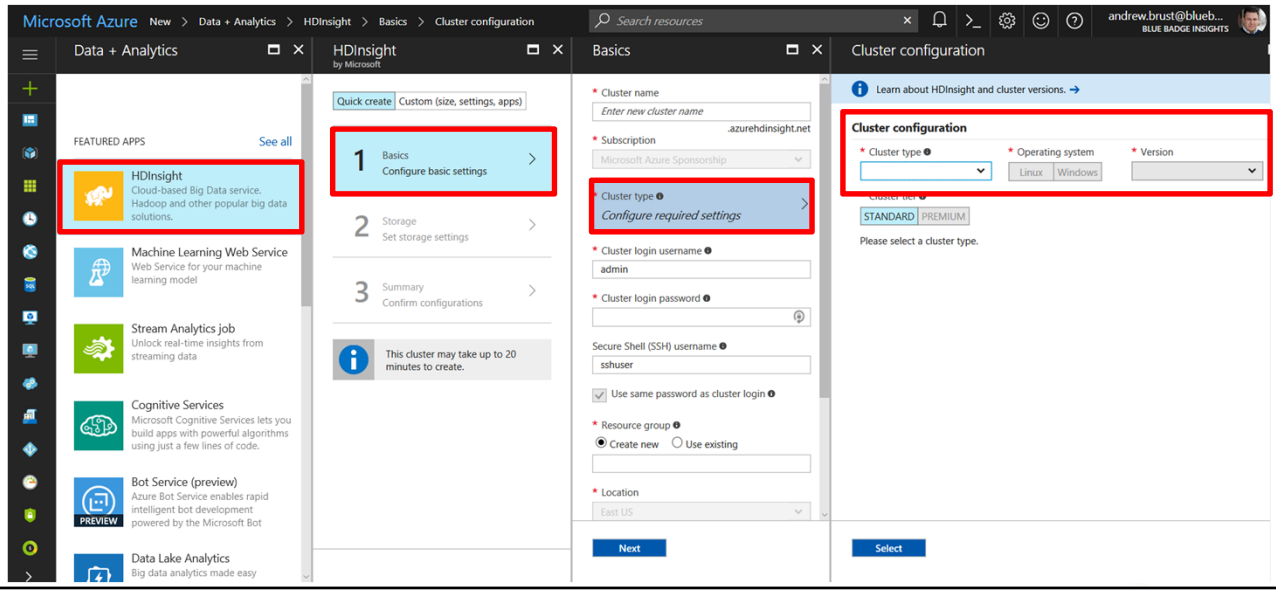

Azure HDInsight Provisioning

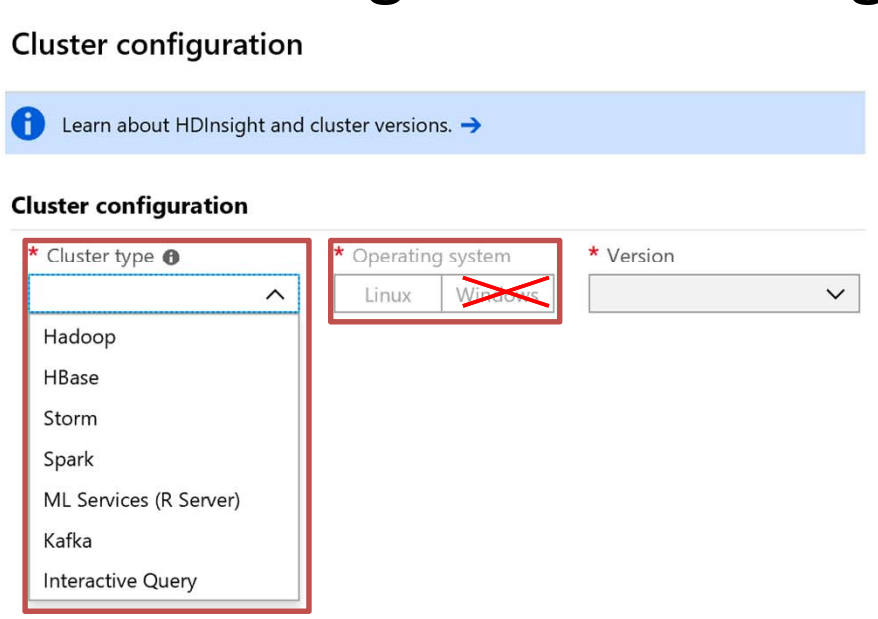

HDInsight Provisioning

demo

Working with HDInsight

- Apache Ambari
  - For Hive queries and cluster monitoring
- Access via PowerShell and HDInsight cmdlets
  - Need to install PowerShell for Microsoft Azure
  - **Run you PowerShell client as administrator**
- SSH into head node
  - Use PuTTY or new SSH client on Windows 10
  - To ***username***@***clustername***-ssh.azurehdinsight.net

Submitting, Running and Monitoring Jobs

- Upload a JAR
- Run at command line (PowerShell or SSH Command line) passing JAR name and params

Clients Options: Command Line via SSH

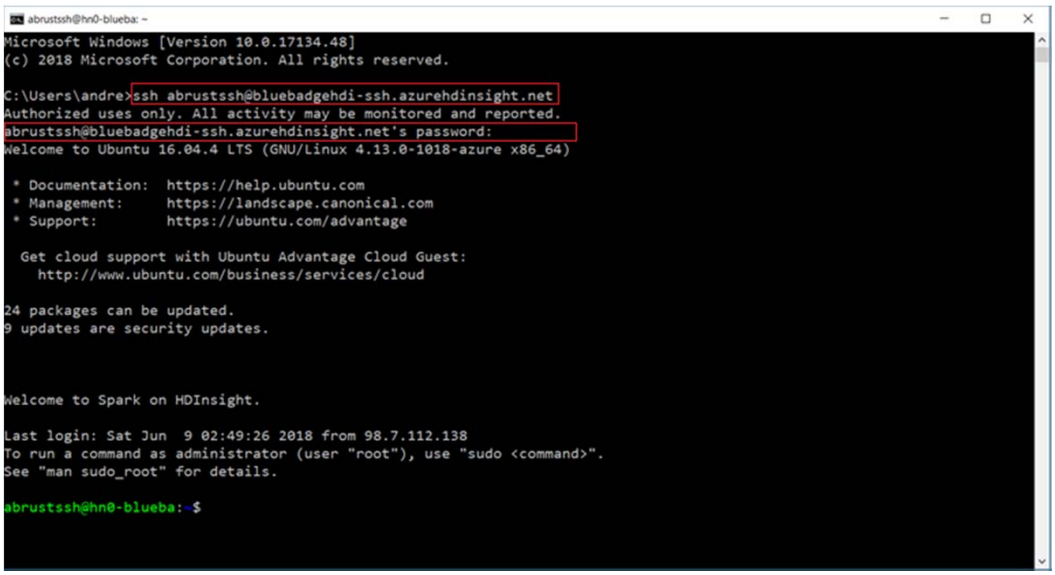

Clients Options: PowerShell

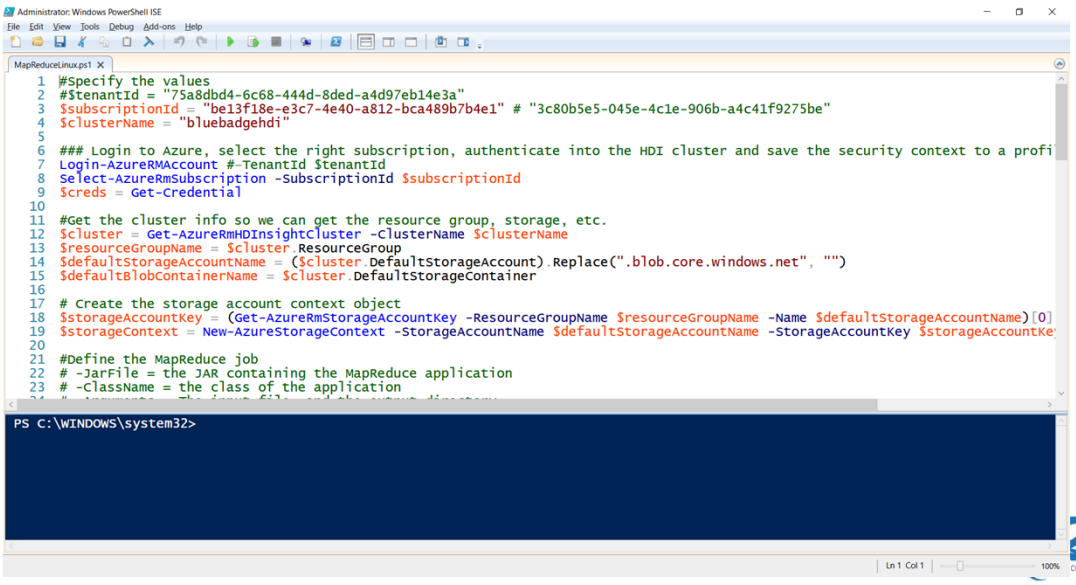

WordCount Code;
Running MapReduce Jobs

demo

Hive

- Originally: a SQL abstraction over MapReduce
- Has evolved to run on Tez
  - Along with Project Stinger, achieved 100x improvement over Hive on MR
- Spark SQL is a cousin of Hive
  - More later
- Hive advanced features
  - External tables
  - UDFs

Clients Options: Ambari

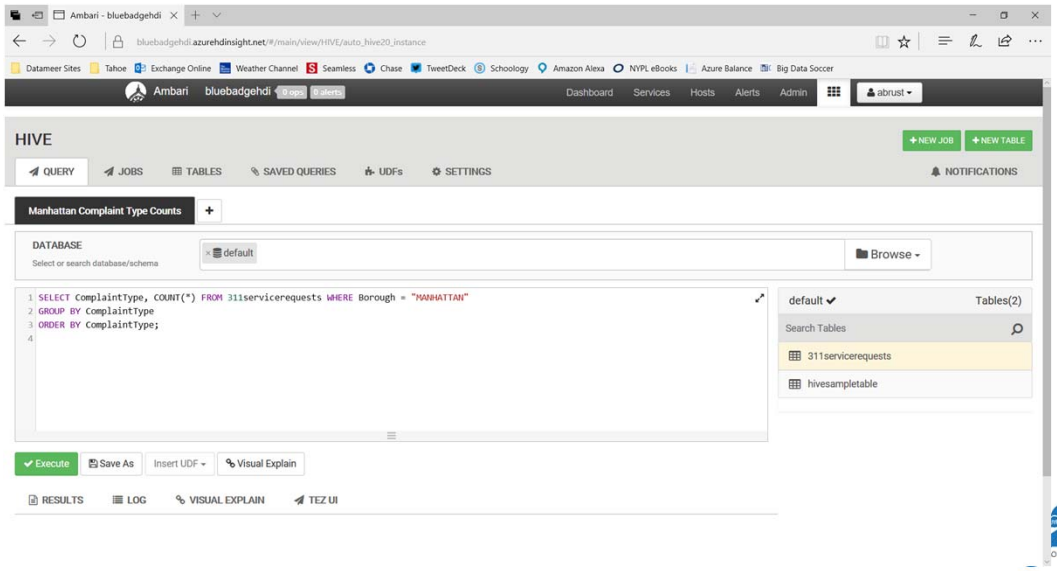

Hive

demo

Hive LLAP, Hive 3

- Stands for “Live Long and Process”
- A Hive-on-Tez variant that uses caching heavily for enhanced performance
- In preview on HDInsight as “Interactive Query” cluster type
- Hive 3: Integrating Apache Druid

Impala, Hive on Spark

- Hive-compatible MPP engine that works directly against HDFS
- Apache Impala was originally a Cloudera project
- Hive-on-Spark is a Cloudera-led enhancement to Hive that has it run on Spark instead of MR or Tez
- Neither one common on non-Cloudera Hadoop clusters

Pig

- Also a programming language abstraction over MapReduce
- Language is called Pig Latin
- Can be used for interactively and for queries
- More often used for data transformation, from scripts

Clients Options: Azure Data Lake Tools in VS

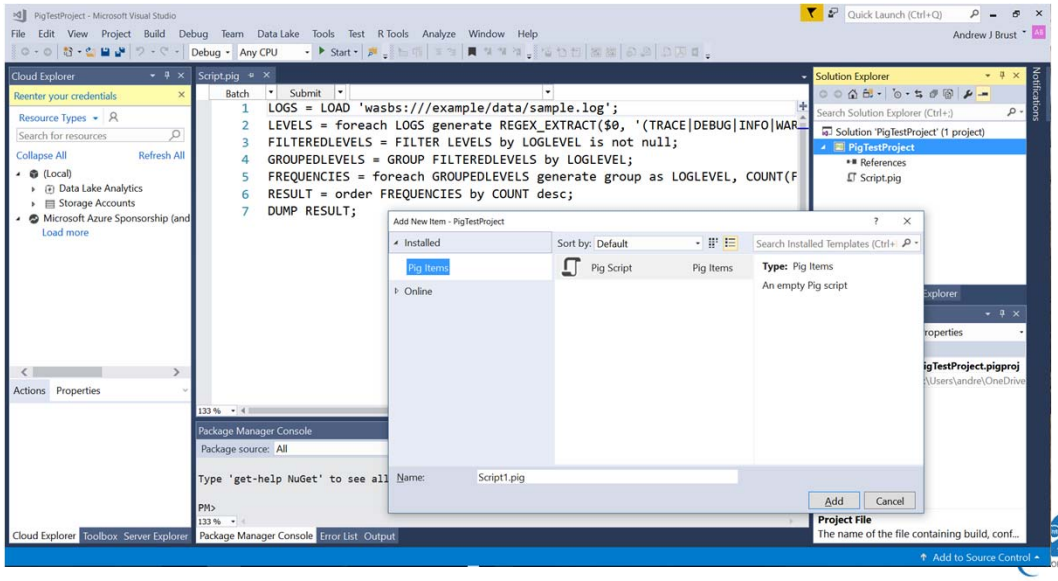

Clients Options: .NET SDK for Hadoop

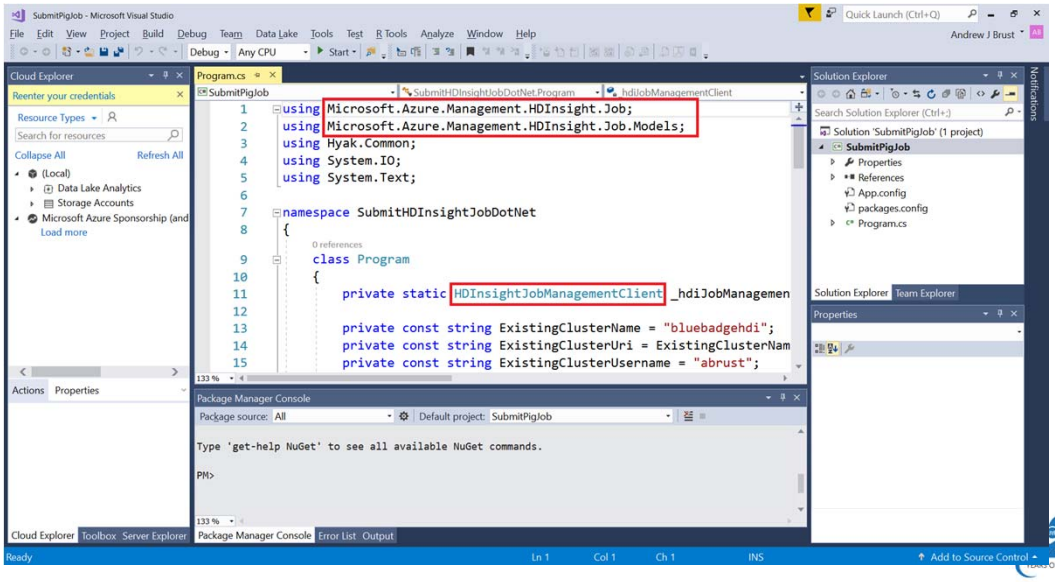

Clients Options: cURL

Pig

demo

AZURE DATA LAKE

Azure Data Lake Store

- Based on Azure BLOB storage, but…
- No file size limits
- Resources added as needed for scale
- WebHDFS compatible
- Certain HDInsight cluster types can use it instead of Blob storage
- Third parties beginning to support
- NEW! ADLS Gen2 (in preview) is a perfect superset of BLOB storage

Azure Data Lake Analytics

- Lets you do big data analytics on data stored in ADLS
- ADLA jobs run on YARN/HADOOP
- Jobs are run on-demand; no dedicated HDInsight cluster involved
- Right now, only job type supported is U-SQL…

U-SQL

- U-SQL
  - Work with flat files or create databases
    - DBs allow for indexing and partitioning
  - Looks like T-SQL, but allows inclusion of C# code...either for inline expressions, or for UDFs
  - Allows batch operations on whole sets of files using wildcard patterns.
  - Not a business user tool, but an *excellent* abstraction layer on Hadoop for developers
- As part of Azure Data Lake Analytics
  - Runs Hadoop jobs behind the scene – but server-less/cluster-less
  - Native storage is Azure Data Lake Store, but can access data in Azure Blob storage too

Azure Data Lake Analytics/U-SQL

demo

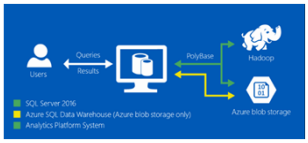

POLYBASE

PolyBase

- A “bridging” technology to connect SQL Server to data in Hadoop or Azure Blob Storage
- Makes the Hadoop data look like SQL Server data via “EXTERNAL” tables
- Query as normal; even join with physical tables
- First appeared in Parallel Data Warehouse/APS and Azure SQL DW
- Now included in SQL Server 2016 Enterprise
- Can create physical table with CREATE TABLE…AS SELECT… (CTAS)

PolyBase

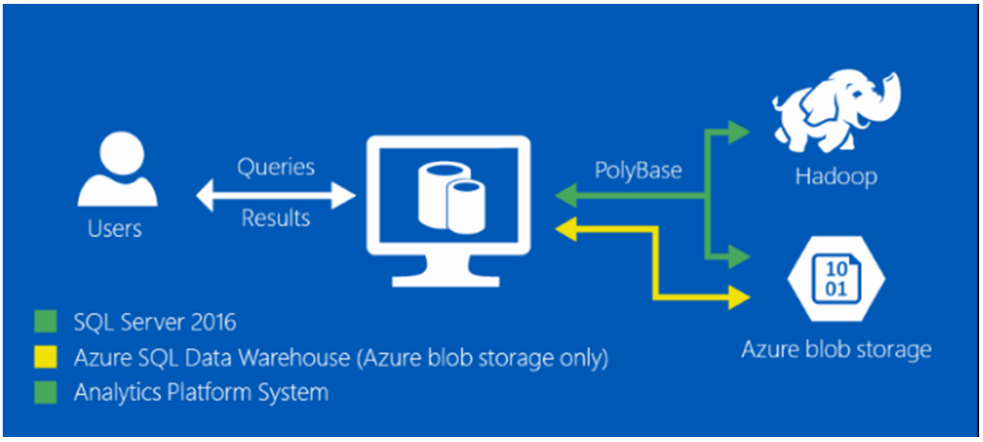

Notes

- Data may be moved and processed by SQL Server's engine and optimizer, or may be "pushed down" to Hadoop, or both
- For DW versions of SQL, query is distributed
- Config can be tricky
- Java install is a prerequisite

Relevant T-SQL

- Prepatory:
  - EXEC sp_configure 'hadoop connectivity', *x*
  - RECONFIGURE;
  - CREATE MASTER KEY ENCRYPTION
- Next:
  - CREATE DATABASE SCOPED CREDENTIAL
  - CREATE EXTERNAL DATA SOURCE
  - CREATE EXTERNAL FILE FORMAT
  - CREATE EXTERNAL TABLE

PolyBase

demo

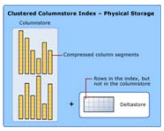

COLUMNSTORE INDEXES

In Analytics…

- Geared to reporting and visualization
  - Read frequently, write seldom
- Table scans are expected
- Aggregation (think GROUP BY) is *de riguer*
- Extensive normalization is bad
- You only care about values in a small set of columns…maybe even just one
  - The rest are used with WHERE and HAVING, to filter
- Tables that track location and time are common

A History of Columnstore Indexes

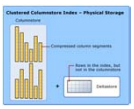

- SQL Server 2012: Nonclustered Columnstore Indexes (NCCIs) added to product
  - Read only
- SQL Server 2014: Clustered Columnstore Indexes (CCIs) added
  - Read/Write
- SQL Server 2016: Numerous enhancements to CCIs

Vector Processing

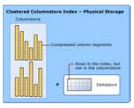

- Intel x86 CPUs have, since supported “single instruction multiple data” (SIMD) operations since the 1990s
- These process data in parallel, handling multiple data points simultaneously
- This is called vector processing
- SQL Server NCCIs and CCIs can take advantage of it

Useful Applications

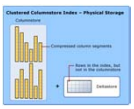

- Data Warehouse/Data Mart scenarios
- In combination with DirectQuery feature in SSAS Tabular and Power BI
- In combination with R Services

Vector Processing and "Batch" Mode

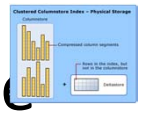

- SQL Server can burst into a vector processing mode
- Instead of iterating through rowsets, one row at a time, it can handle rows in batches
- So it's called "batch mode" and it's *fast*
  - (Not to be confused with batch processing, which can be slow)

Sanity Check

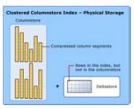

Columnstore Index Scan (NonClustered)

Scan a columnstore index, entirely or only a range.

Physical Operation	Columnstore Index Scan
Logical Operation	Index Scan
Actual Execution Mode	Batch
Estimated Execution Mode	Batch
Storage	ColumnStore

Making Sure it Works

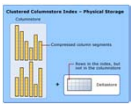

- CIs are fastest when Batch mode kicks in
  - Difference can be negligible otherwise
  - Check Query Plan to make sure
- And meet the prerequisites…

Prerequisites

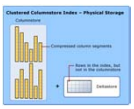

1. More than one CPU core	2. Maximum Degree of Parallelism (MDOP) set to 0	*Lots* of data
• (Careful on those VMs!)	• Or a value between 2 and 64, if you want to limit it • Use SSMS server properties sheet or sp_configure and RECONFIGURE	• Millions of rows, or don't bother

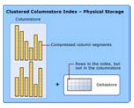

Columnstore Indexes

demo

Distributed, In-Memory
Big Data Platform

APACHE SPARK

Spark

- Wildly popular open source project, focuses on distributed in-memory processing versus on-disk
- Can use it independently of Hadoop, but most people use it with Hadoop/HDFS
- Very popular component: Spark SQL
  - Allows HiveQL queries against Spark (Power BI can use this)
- Also: Spark Streaming, MLlib, GraphX
- Spark now supported on HDInsight

Spark on HDInsight

- HDInsight Spark clusters include the Jupyter and Zeppelin “notebook” user interface
- Allow interspersal of text, code, and code output, including visualizations
- Supports Python (PySpark) and Scala
- Includes *very* helpful tutorial notebooks

Jupyter Notebooks

- Notebooks combine code, text and data visualization capabilities
- Text and code “cells” are interspersed. Code can be executed in place.
- Jupyter originally called iPython and hosted only Python code; now hosts numerous langauges
- On HDInsight, Jupyter Notebooks can host Python, Scala and R code, running against Spark
- See also: Azure Notebooks

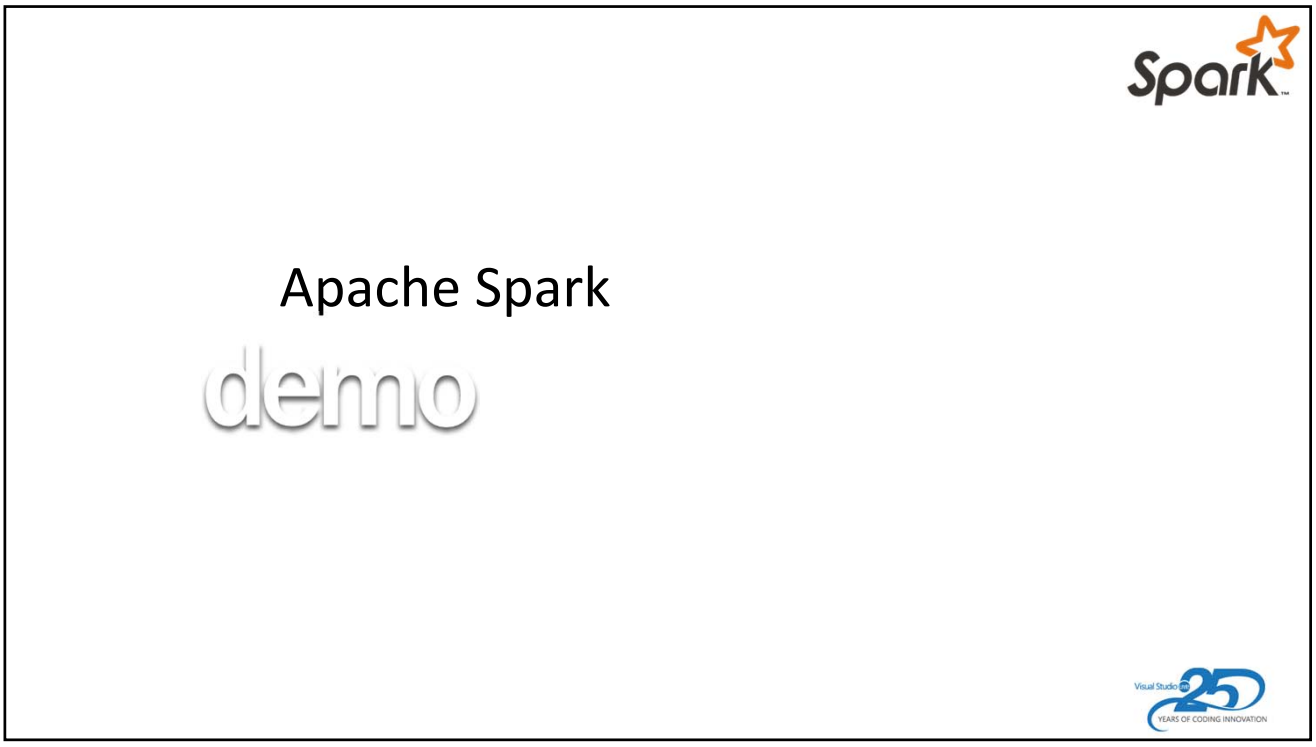
Spark
Apache Spark
demo
Visual Studio 25 YEARS OF CODING INNOVATION

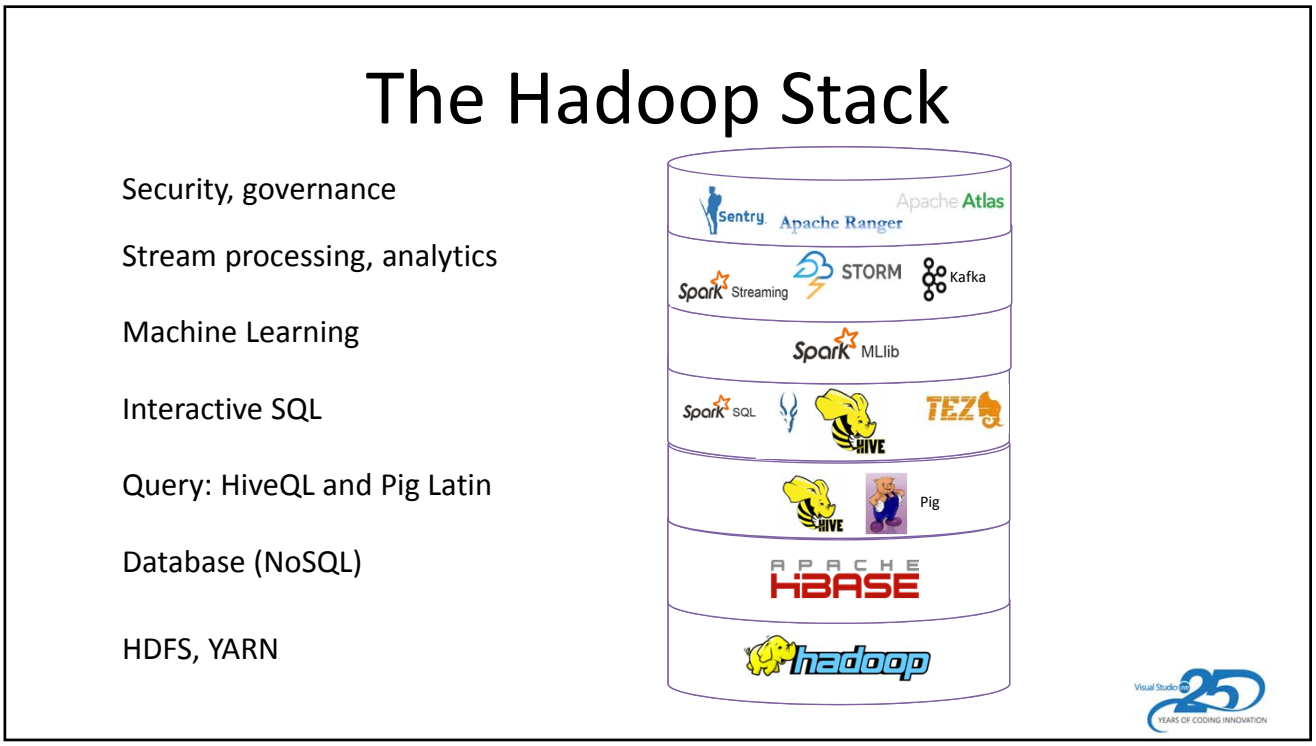
The Hadoop Stack
Security, governance
Stream processing, analytics
Machine Learning
Interactive SQL
Query: HiveQL and Pig Latin
Database (NoSQL)
HDFS, YARN
Sentry
Apache Ranger
Apache Atlas
Spark Streaming
STORM
Kafka
Spark MLlib
Spark SQL
HIVE
TEZ
HIVE
Pig
APACHE HBASE
hadoop
Visual Studio 25 YEARS OF CODING INNOVATION

Microsoft's Modern BI Platform

POWER BI

Power BI

- Based on same columnar, in-memory BI engine as SQL Server Analysis Services Tabular mode
- Free Desktop and Mobile apps
- For individual users, 2 cloud subscription levels: Basic (free) and Pro ($10/month/seat)
- Easy to use, extensible, embeddable, connects to a huge array of conventional and cloud data sources
  - Growing DirectQuery support
- Highly integrated across Microsoft stack

Power BI Ingredients

Power BI Desktop

- Windows desktop app
- Acquire, shape data query editor
- Visualize data with report view

Browser Environment

- www.powerbi.com
- Edit, consume
- On-prem gateways

iOS, Android, Widows Universal Apps

- iPad, iPhone, Android phones and tablets
- Windows tablets, PCs
- Consumption only

Power BI Desktop

- Windows Desktop Application
- Has a "main window," akin to the Excel Power View Add-In, for report authoring and some data modeling
  - Report view
  - Data view
  - Relationships view
- Has a Query Editor window, akin to the Excel Power Query Add-In, for data import and transformation
- Can save files (.pbix) locally and an publish them to powerbi.com

Power BI Query Editor: Overview

- Launched with Get Data option (from ribbon or splash page)
- Re-entered using Edit Queries ribbon button
- Use it to import and shape data
- Use Close & Load ribbon button when done
- Try not to confuse this window with the data view in the main window

Get Data, Query Editor

demo

Power BI Reports Overview

- Data exploration and visualization client
- Visualizations work as filters, too
- Design and view experiences are unified

On-Premises Gateway

- Permits import and scheduled refresh of on-prem data in cloud copy of report
- Personal mode:
  - Runs as app for single user
- Enterprise mode:
  - Runs as service for multiple users
  - “DirectQuery” supported for numerous data sources
  - “Live Connection” supported for SSAS (Tabular or MD)
  - Supports PowerApps, Azure Logic Apps, Microsoft Flow and Azure Analysis Services (preview)

The Views

- Report: the report designer/viewer
- Data: where you can model the data
  - Rename/delete/hide columns and tables
  - Sort by a column (ascending or descending)
  - Add DAX measures and calculated columns
  - Set data types and categories
- Relationships – Where you can view and edit relationships
  - But you must create them with the Manage Relationships dialog

Power BI Reports

demo

Power BI Cloud Service

- Authoring and consumption tool
- Can create three things
  - Dataset
  - Report
  - Dashboard
- Publish report from PBI Desktop, get link to cloud version
- Also available: “Quick Insights”

Dashboards

- A collection of “pinned” visualizations from existing Power View reports
- Pin entire reports, too!
  - Single visualizations are not interactive
- What you can pin:
  - Web content, images, video, text boxes
  - Visualizations from Quick Insights
  - Excel spreadsheet assets
  - SQL Server Reporting Services assets
  - Camera photos (via iPhone App)

Power BI Service

demo

Q&A

- Natural language query interface to data in underlying model
- Available at top of dashboard
  - Now available in reports too
  - And as authoring tool
- Generates visualization as you type
- Visualization is pin-able

Power BI Premium

- New subscription level for Enterprise use:
  - Unlimited consumption users; Professional subscription still required for each authoring user
  - Dedicated infrastructure; paid for by the number and type of server nodes
  - Starts at $4,995/month for P1 node with 8 cores, 25GB RAM
- Includes on-premises capabilities:
  - Power BI Report Server: Actually a superset of SQL Server Reporting Services. (Available w/o power BI subscription for SQL EE+SA customers.)
  - Licensed for same number of cores included in cloud subscription
  - Reports only; no dashboards

Where Microsoft BI Started…

ANALYSIS SERVICES

Data Migration

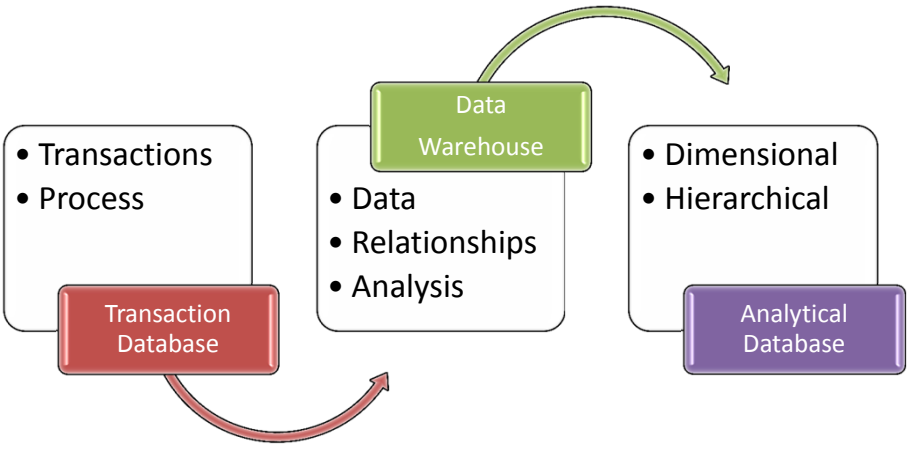

SQL Server Analysis Services

- Built for analysis
- Included with SQL Server Standard, Enterprise
- And you can use the Microsoft stack that you know and love

From Data Warehouse to OLAP

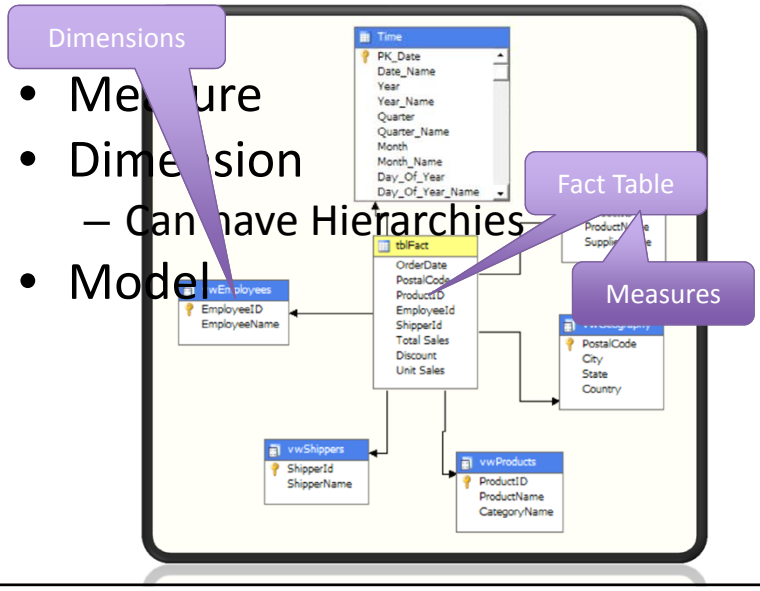

- Me ure
- Dim sion
  - Can have Hierarchies
- Model

Analysis Services Modes

- **Multidimensional** or **Tabular**
- Tabular is newer, same tech as Excel/PowerPivot data models and Power BI
- Lots of investment in Tabular in SSAS 2016
- We'll look at Tabular today

Analysis Services Tabular Mode

- SSAS Tabular Mode uses a columnar storage engine in place of a multidimensional one
- Must choose mode for SSAS instance at install time
- Can have default instance with one, named instance with the other
- Can create an SSAS Tabular database project by importing an Excel workbook with PowerPivot model
- SSAS tabular models support partitions, roles, translations, display folders

Calculated Columns and DAX

- Formula-based columns may be created
- Formula syntax is called DAX (Data Analysis eXpressions).
  - Not to be confused with MDX or DMX. Or DACs.
- DAX expressions are similar to Excel formulas
  - Work with tables and columns; similar to, but distinct from, worksheets and their columns (and rows)
- =FUNC('table name'[column name])
- =FUNCX('table name', <filter expression>)
- FILTER(Resellers,[ProductLine] = "Mountain")
- RELATED(Products[EnglishProductName])
- DAX expressions can be heavily nested

Analysis Services

demo

Azure Analysis Services

- In preview now
- Platform as a Service offering for Analsis Services Tabular
- Supports Analysis Services 2017 features
- Compatible with Excel, Power BI
- Can use Visual Studio Analysis Services Projects tooling or new browser based tools
- Can use same on-prem gateway as Power BI for refresh of models from on-prem data sources

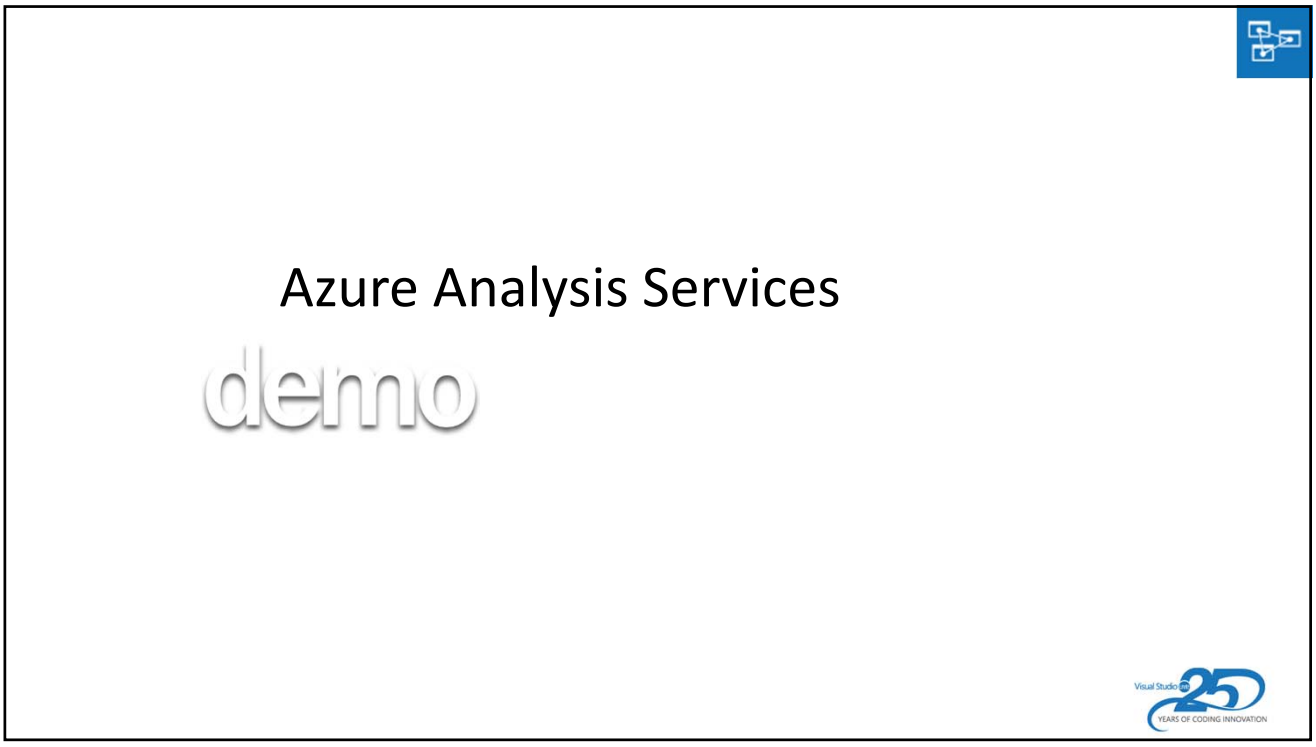
Azure Analysis Services
demo


CLOSING THOUGHTS

How Do BI and Big Data Relate?

- At the root of each lies the idea of grouping and aggregating
- The Reduce step in MapReduce is all about that
- On the DW/BI side, so is defining dimensions and drilling down by them
- And there is a pretty strong linkage between dimensions/reducer groupings on the one hand and machine learning features on the other
- Think of it this way…

Connect the Dots

1	2	3	4	5
SQL: SELECT	Hadoop: Map	BI: Measure	Visualization: Values	PivotTable: Data Cells

Connect the Dots

1	2	3	4	5
SQL: GROUP BY	Hadoop: Reduce	BI: Dimension	Visualization: Axis	PivotTable: Column or Row

Integration Matrix

	SQL Server RDBMS	HDInsight	Power BI	Analysis Services	Excel	Reporting Services	R	U-SQL	.NET
SQL Server RDBMS		●	●	●	●	●	●		●
HDInsight	●		●	●	●		●	●	●
Power BI	●	●		●	●		●		●
Analysis Services	●	●	●		●	●			●
Excel	●	●	●	●					●
Reporting Services	●	●		●			●		●
R	●	●	●					●	
U-SQL		●					●		●
.NET	●	●	●	●	●	●		●	